

Business Brain

Five-minute demo video: script and production guide

Every figure quoted is a live value from the seeded database.

What you are making

A five-minute screen recording with voice-over that walks a judge through the product: the world model, the money view, the ranked attention queue, an agent decision approved live on camera, counterfactual simulation, the admin analytics, and a self-test suite run in real time.

Runtime	5:00. The narration is 722 words, about 4:50 of talking at a normal pace, which leaves roughly ten seconds of breathing room for the clicks and the two moments where you wait for the app to respond.
Method	Two passes. Record the screen silently first, then record voice to picture, then combine.
Tools	QuickTime Player and iMovie. Both already on your Mac, both free.
Hard rule	Record in a browser window at least 1440 px wide. Below 900 px the app switches to its mobile layout and the demo looks like a phone app.

The two moments that sell it

At **2:30** the system shows what it *rejected* and why. At **4:25** twenty-one assertions run against live data on camera in about a second. Most demos are a tour of screens; these two are not. Do not rush either.

Before you record

Do this once, about ten minutes before the take.

1. Confirm all four services are up

Easiest check: open `/app/admin/system` in the browser. The first tile reads **Dependencies healthy 2 / 2** when Mongo and Redis are both live, and the page only loads at all if the API and the web app are running. If you would rather check from the terminal, this is one line and safe to paste:

```
cd /Users/dhruvdawar11/Desktop/Projects/Razorpay-Build-Hackathon

.venv/bin/python -c "import socket;print(' | '.join(n+' '+'UP' if socket.socket().connect_ex(('127.0.0.1',p))==0 else 'DOWN') for n,p in [('mongo',27017),('redis',6379),('api',8000),('web',3000)]))"
```

If the API or the web app is down, restart it:

```
cd apps/api && SEED_ON_START=true ../../.venv/bin/python -m uvicorn app.main:app --port 8000 &

cd apps/web && NEXT_PUBLIC_API_URL=http://localhost:8000 npm run dev
```

2. Set the stage

- 1 **Sign in** at `http://localhost:3000` with `merchant@atlas.local` and `atlas-demo-change-me`.
- 2 **Dark mode on** using the toggle at the bottom of the sidebar. It reads far better on video.
- 3 **Maximise the browser** so the window is at least 1440 px wide. Hide the bookmarks bar with `Cmd-Shift-B`. Close every other tab.
- 4 **Turn on Do Not Disturb** so no notification lands mid-take.
- 5 **Leave the pending approval alone.** Exactly one action is waiting for approval and you approve it live at 3:05. If you click it in rehearsal you lose the best moment in the demo.
- 6 **Do not run a brain cycle** before recording, or the numbers below will drift off the script.

3. Pre-open these eight tabs, in this order

Then you never fumble a URL on camera. Switch with Cmd-1 through Cmd-8.

TAB	URL	SCREEN
1	/app	Overview
2	/app/money	Unresolved money
3	/app/attention	Attention
4	/app/decisions	Decisions
5	/app/simulate	What if
6	/app/admin/analytics	Analytics
7	/app/admin/agents	Agents and trust
8	/app/admin/testing	Test lab

4. Record in two passes

- 1 Pass one, screen only.** QuickTime Player, File, New Screen Recording. Set Microphone to None. Follow the shot list slowly, pausing two to three seconds on each screen. You will be talking over this later, and rushed clicking is the single thing that makes a demo feel amateur. Aim for about 5:15 so you have trim room.
- 2 Pass two, voice only.** QuickTime Player, File, New Audio Recording. Play the screen recording on one side and this script on the other, and narrate to picture. Record it in the eight sections below, not in one take. Fluff a line and you redo one section instead of restarting.
- 3 Combine.** iMovie, new project, drag in the screen recording, drop each audio section underneath and nudge it into place. Mute the original video audio. Share, File, 1080p, High quality.

If you would rather do one take, set QuickTime's Microphone to your mic and narrate live. Faster, but budget for retakes.

The script

Blue blocks are what you **say**. Orange lines are what you **click**. Timings are cues, not handcuffs.

0:00

Hook

Tab 1, Overview. Sit still, do not scroll yet.

A merchant opens ten dashboards to answer one question: *am I okay?* Payments in one place, settlements in another, payouts in a third. Nobody joins them up.

This is Business Brain. One world model per merchant. Events go in, and a single state vector comes out that every agent reads from. Here is Atlas Electronics, live.

0:30

The world model

Tab 1. Scroll slowly to the forecast, then the state vector.

Revenue in the last twenty-four hours, thirteen lakh, down six percent against its own seven-day baseline. Cash twenty-five lakh, projected to twenty-seven ninety-two, against a reserve floor of twenty lakh the merchant set themselves.

But the label says **critical**. Not because cash is short, cash is fine. Because the failure rate is fourteen point six percent against a five point nine baseline. Two and a half times normal, and the system caught it before the merchant did.

All of it is computed, not stored. Below: the cash forecast with a p10 to p90 band, three days of financial weather, and the full state vector the agents read before they propose anything.

ON SCREEN Hover one point on the forecast band so the tooltip appears.

1:15

Where is the money

Tab 2, Unresolved money.

Second question every merchant has: where actually *is* my money?

Thirty-six lakh has left a customer but is not usable cash yet, split by how recoverable it is. Twenty lakh comes back inside forty-eight hours. Pending settlements, twenty-one lakh, ninety percent recoverable. Scheduled payouts, twelve lakh, and that is committed, it is leaving.

Not guesses. Each recovery rate is this merchant's own historical rate, and every chart has a table view, so no number is trapped behind a hover.

ON SCREEN Click **Table** on the first card for two seconds, then click back to **Chart**.

1:50 Attention and causality

Tab 3, Attention. Scroll to the causal graph.

So what do I do? The attention queue ranks by impact times urgency times confidence times reversibility. A real score, not a rules list. Two items: settlements delayed, ten point four lakh, and payment failures rising, marked high, with the reasoning attached.

Every item carries its evidence IDs, the exact fields it came from. And here is the causal graph: edges whose strength is measured from this merchant's own time series, not assumed. Failure rate against success rate, minus one. Measured, not hardcoded.

ON SCREEN Point the cursor at the failure_rate to success_rate row showing **-1.000**.

2:30 The decision, and a live approval

Tab 4, Decisions. This is the centrepiece.

Now the part that matters. Four specialist agents bid on this. The orchestrator scored them under the merchant's own objective weights and picked one: retry the failed payments, thirteen point four lakh.

But here is what I care about. It shows you what it *rejected*. Do nothing was considered and dropped, because doing nothing leaves reserve-breach probability at eight point six percent.

On the right, the evidence it rests on, copied from world state at decision time, and the expected impact, simulated before anything executes.

And it cannot act on its own. This is a financial action, so policy holds it. Let me approve it.

ON SCREEN Click the **Approvals** tab, then **Approve** on the pending row. Wait for the counters to move before you speak again.

Approved, executed, verified, and the measured impact went into the outcome ledger. Not the estimate. What actually happened.

3:20

What if

Tab 5, What if.

Before you commit, you can ask what if. Same engine the orchestrator scores its own options with.

ON SCREEN Click **Compare options**. Wait for the six rows to appear.

Six scenarios from the same starting state. Delay payouts twelve hours, plus three lakh. Everything at once, plus six and a half lakh, and breach risk drops from eight point six to seven percent. Nothing here touches money. It only scores.

3:50

Analytics and earned trust

Tabs 6 then 7.

Everything the brain sees, the admin sees.

ON SCREEN Tab 6. Scroll through the charts while you talk.

Thirty days: two point nine eight crore captured at ninety-five point seven percent success. UPI leads volume and also has the worst failure rate. The full funnel. And outcome calibration, predicted against measured, so you can see where the agents are wrong.

ON SCREEN Switch to Tab 7.

Trust is earned, per merchant. Four agents, each declaring its capabilities. Trust moves only on a verified outcome, every bid is logged win or lose, and circuit breakers cap how often any agent can act.

4:25

The self-test, live

Tab 8, Test lab. Do not skip this.

Last thing, and this is the part I would want to see if I were judging.

ON SCREEN Tick **Include write checks**, then click **Run self-test**. It takes about a second and a half.

Twenty-one checks, end to end, against live data. Not mocks. That forecast bands are ordered, that unresolved money reconciles to its buckets, that the kernel rejects cross-merchant calls, that replaying an event ID is deduplicated, and that an action's lifecycle lands in the outcome ledger exactly once.

All twenty-one green, in about a second. This suite found two real bugs during development: a double-counted outcome on replay, and an inverted sign in the causal graph.

4:50

Close

Stay on the green result.

One world model. Agents that bid and explain what they rejected. A safety kernel that will not let them move money unattended. Every layer visible and testable from the console. That is Business Brain. Thanks for watching.

Saying the numbers

Indian numbering trips people up on camera. Read the right-hand column, not the digits.

ON SCREEN	SAY THIS	WHAT IT IS
₹13.13L	thirteen lakh thirteen	Revenue, last 24h
₹25.00L	twenty-five lakh	Cash on hand
₹27.92L	twenty-seven ninety-two	Projected cash, 24h
₹20.00L	twenty lakh	Reserve floor
14.6% / 5.9%	fourteen point six against five point nine	Failure rate now vs baseline
₹36.36L	thirty-six lakh	Total unresolved money
₹20.08L	twenty lakh recoverable	Recoverable inside 48h
₹10.43L	ten point four lakh	Settlements delayed
₹13.39L	thirteen point four lakh	The recommended retry
₹2.98Cr	two point nine eight crore	GMV captured, 30 days
₹1.46Cr	one point four six crore	UPI volume
21/21	twenty-one of twenty-one	Self-test result

If something goes wrong on camera

A page looks empty	Click Refresh in the page header and keep talking. Do not stop the take.
The approval is already gone	Test lab, Action probe tab. Create a refund.create for 9000000 and approve or reject that instead. It is the same pipeline and the same policy gate.
Numbers differ from this script	Read what is on screen, not what is printed here. The story works with any values.
The layout looks like a phone	Your window is under 900 px wide. Maximise it and reload.
The copilot says model unavailable	That is expected without a Gemini key, and it is not an error. It falls back to deterministic templates over the same evidence. The demo does not visit that screen.

One last thing

Do not apologise for anything on camera, and do not narrate what the viewer can already read. Say what the number *means*. The label reading critical while cash looks healthy is not a wart to hide: it is the clearest proof in the whole demo that the state vector is multi-dimensional rather than one number with a colour. Call it out.